

CS6650 Final Project


Team QuickHorse (AKA 快速牛马)




Shibo Zheng
Ruiyi Li



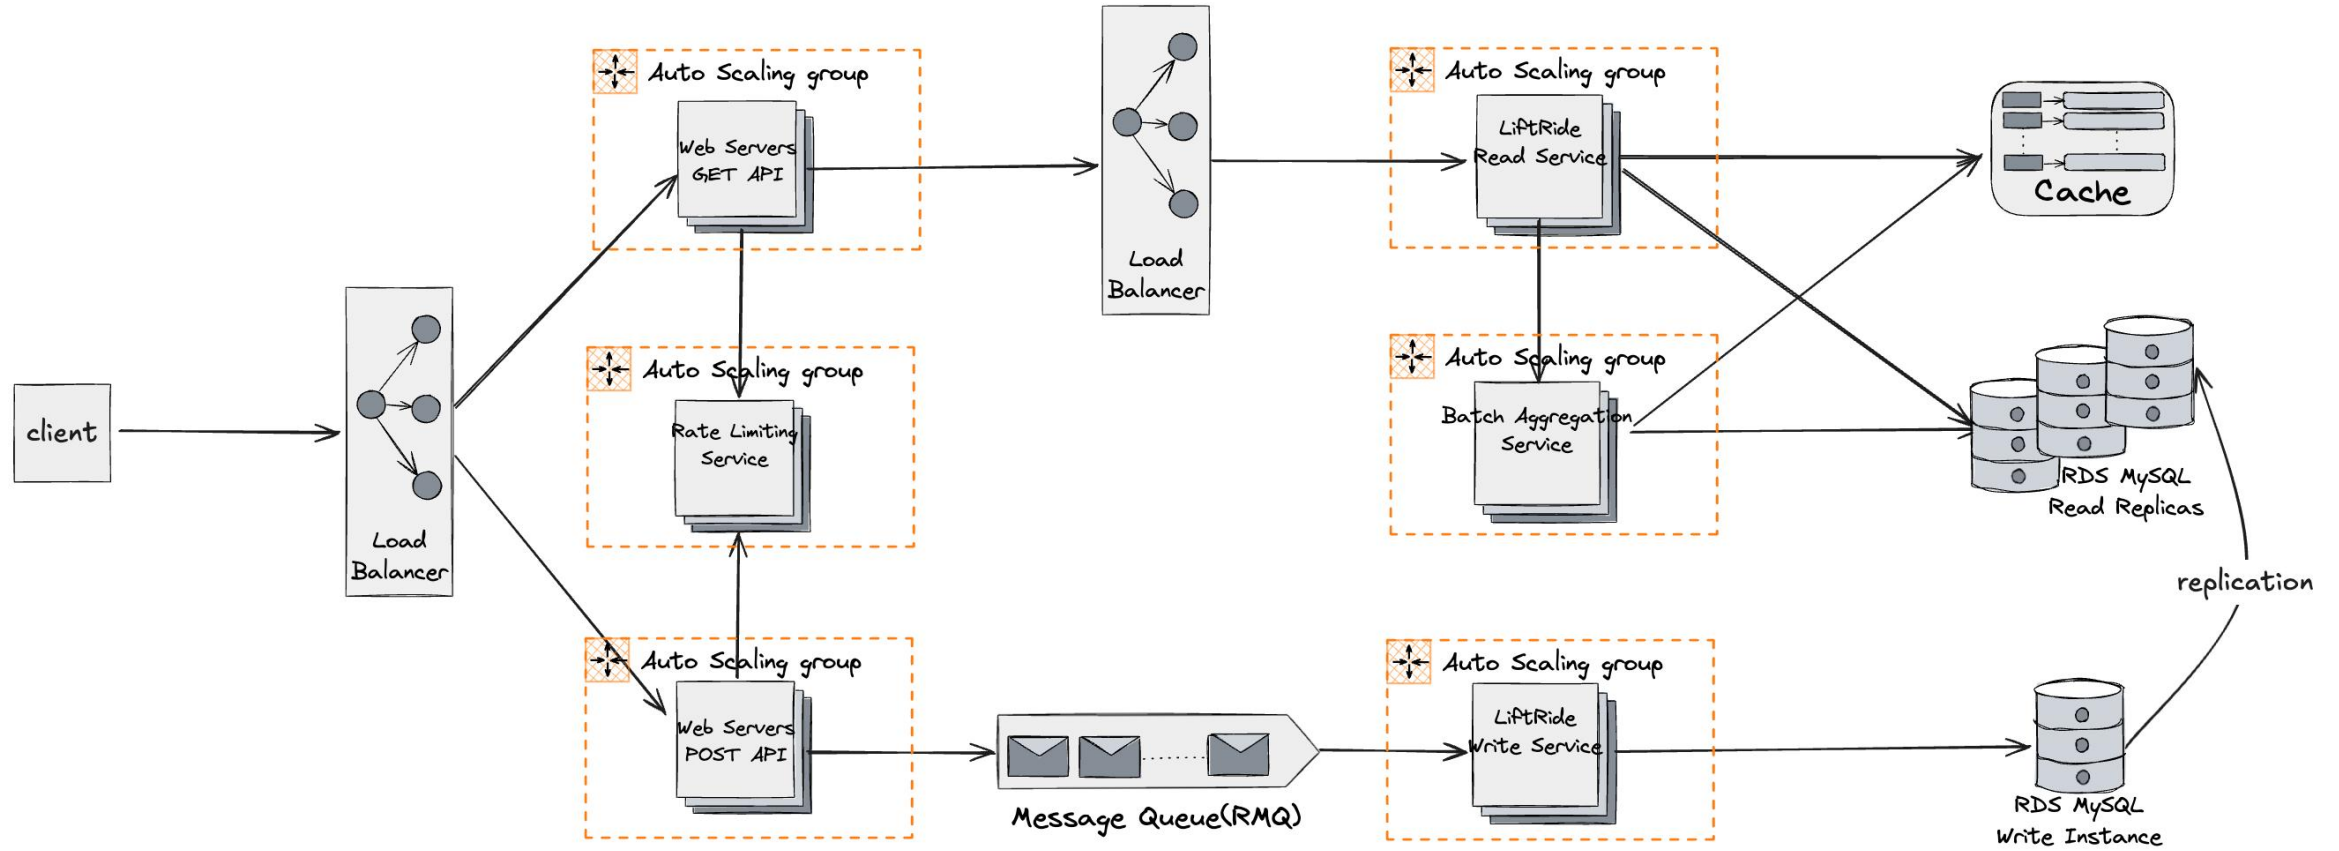
Zhongzhe Liu
Han Yang



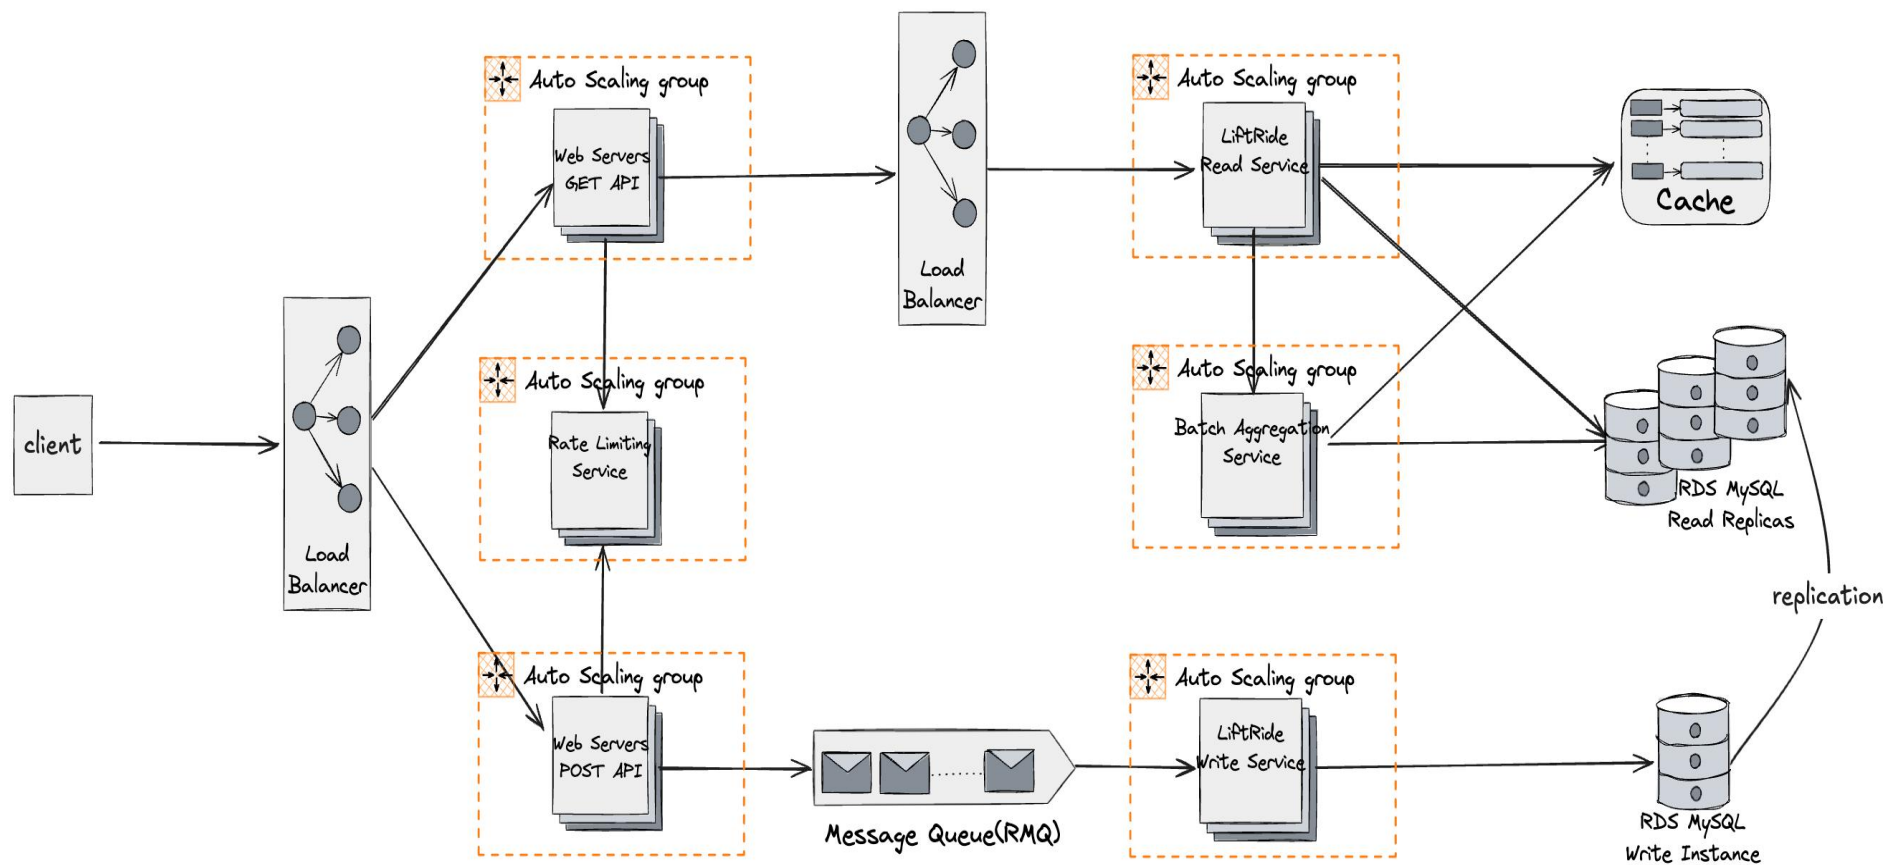
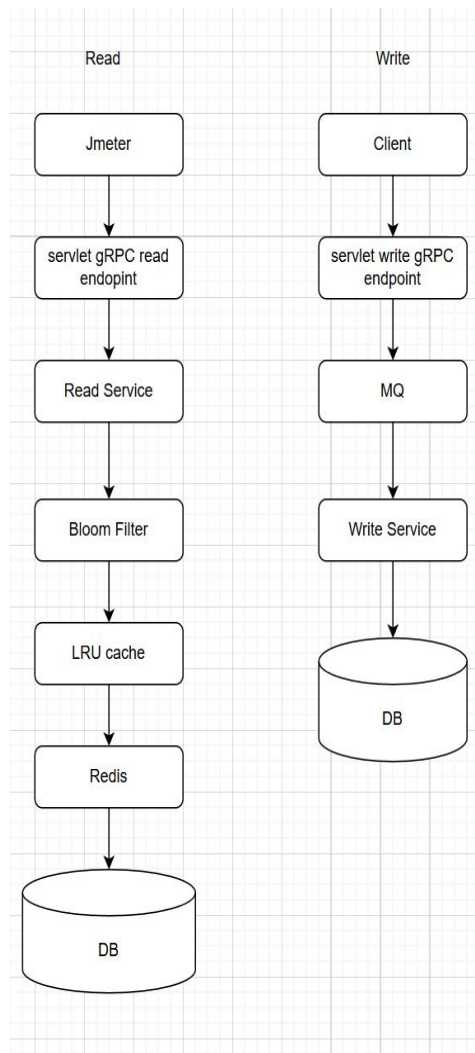
Goal



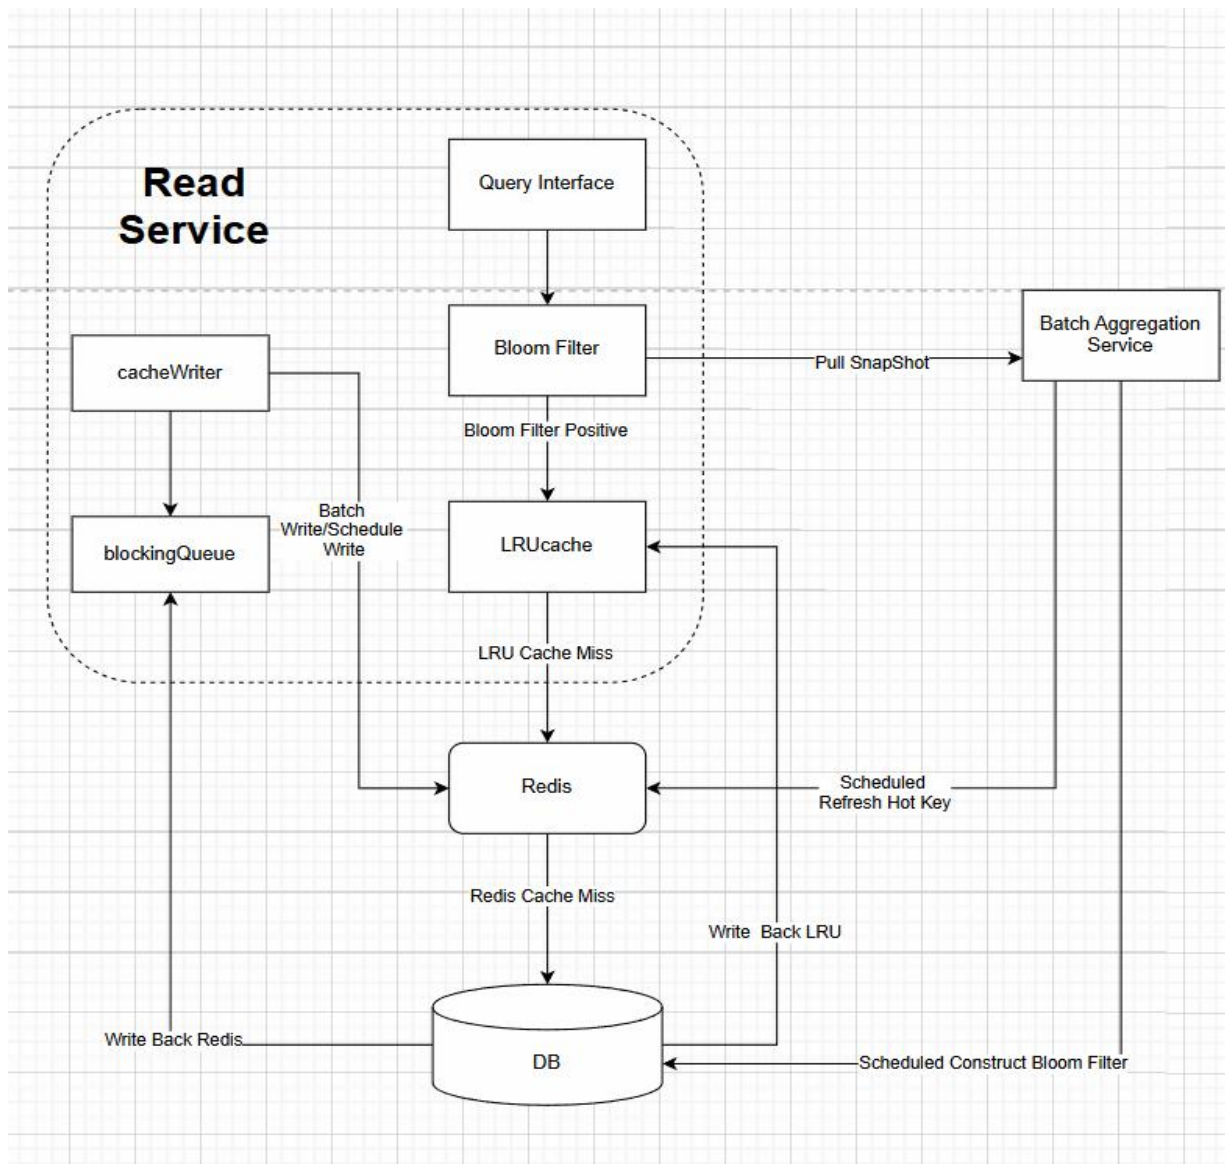
System Architecture



Work Flow



Read Path



- Batch Aggregation Service:
 - Periodically aggregate data from DB
 - Based on strategy:
 - FULL → update Redis + Bloom
 - BLOOM_ONLY → update Bloom only
 - REFRESH_EXISTING_CACHE → update Redis only
 - Builds Guava Bloom Filters, exposes snapshot via gRPC

Microservice Component

1 Client & Jmeter

2 Write Servlet

3 Read Servlet

4 Ratelimiter

5 ServerWriteService

6 DB read service

7 Cache read service

8 Batch aggregation Service

Deployment on AWS

t2.large
Servlet-write

01

t2.large
Write service

03

t2.large
DB instance

05

t2.medium
Redis

07

t2.medium
Rate limiter

09

02

t2.large
Servlet-read

04


t2.large
Read service

06


t2.medium
RabbitMQ

08

t2.large
Batch Aggregation Service



Data Model Design



GET/resorts/{resortID}/seasons/{seasonID}/day/{dayID}/skiers

Get Number of unique skiers at resort/season/day

GET/skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID}

Get total vertical for the skier for the specified sku day

GET/skiers/{skierID}/vertical

Get the total vertical for skier of specified resort, if not specified, return all season

Data Model Design

- Request Body: LiftRide (day, liftID), resort_id, season_id, day_id, skier_id
 1. Each resort can store multiple seasons data
 2. Each resort has multiple lifts
 3. Each skier can go to multiple resorts in any days in any seasons and take mutiple lifts
- Mapping Relation between entitles:
 1. Resort : Season \rightarrow 1 to N
 2. Resort : Lift \rightarrow 1 to N
 3. Skier: LiftRide \rightarrow 1 to N
 4. LiftRide : Lift \rightarrow N to 1

Trade Off between SQL and NO SQL

- Strong Data Integrity with Relationships.
 - Data model includes multiple interrelated entities.
 - Referential integrity is enforced through foreign key constraints.
 - NoSQL DB do not support joins or relational constraints, leading to data duplication or inconsistent references.
- Complex Querying Requirements (Aggregation)
 - Read operation required grouping, filtering, and joining across multiple tables

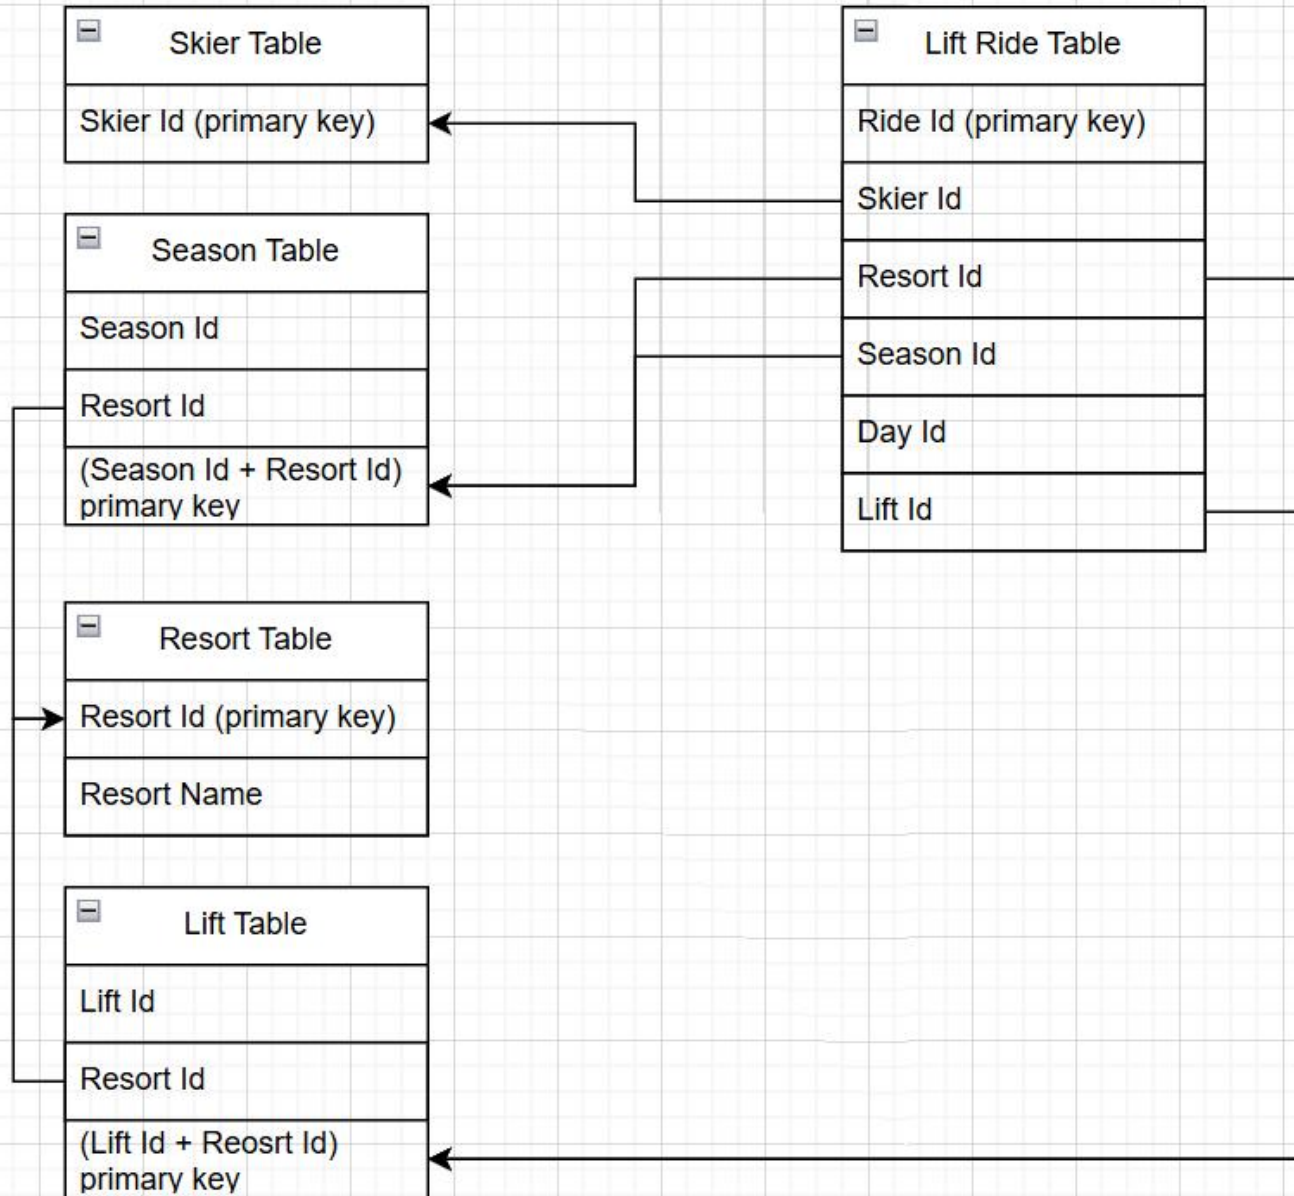
Trade Off between SQL and NO SQL



NOSQL database can
do complex query by
assigning partion key
+ sort key

But, not Scalable, not
Efficient

Redundant data still
exists




Database Schema

Hybrid Solution



- **Multi-tierd cache strategy**
- In Memory LRU cache in read-service
- Redis as second layer cache service
- We also implmented a bloom filter to prevent “Invalid” request
- Why Bloom Filter:
 1. Guarantee no false negatives, but allows false positives.
 2. Light weight data structure
 3. Fast look up
 4. Prevent waste on computing & I/O



Some Optimization

- Decouple bloom filters with Redis
- **Before:** Each ReadService queried Bloom filters via RedisBloom
 - Bloom filters were initialized and stored in Redis
 - We do not want to too much access of Redis!
- **After:** Remove bloom filter from Redis, build and managed locally in memory
 - ReadServices periodically pull updated Bloom filter snapshots from BatchAggregation Service
 - ReadServices perform in-memory mightContain() checks without Redis calls
- **Why**
 - No Redis-Bloom Filter dependency → improved system stability
 - ReadServices become more fault-tolerant and loosely coupled

Some Optimization

- ReadService writes to Redis synchronously after DB fallback

Before:

- Async Read-Through & Cache- Write
- Redis writes are pushed to buffer
- Flushed asynchronously by a dedicated CacheWriter thread

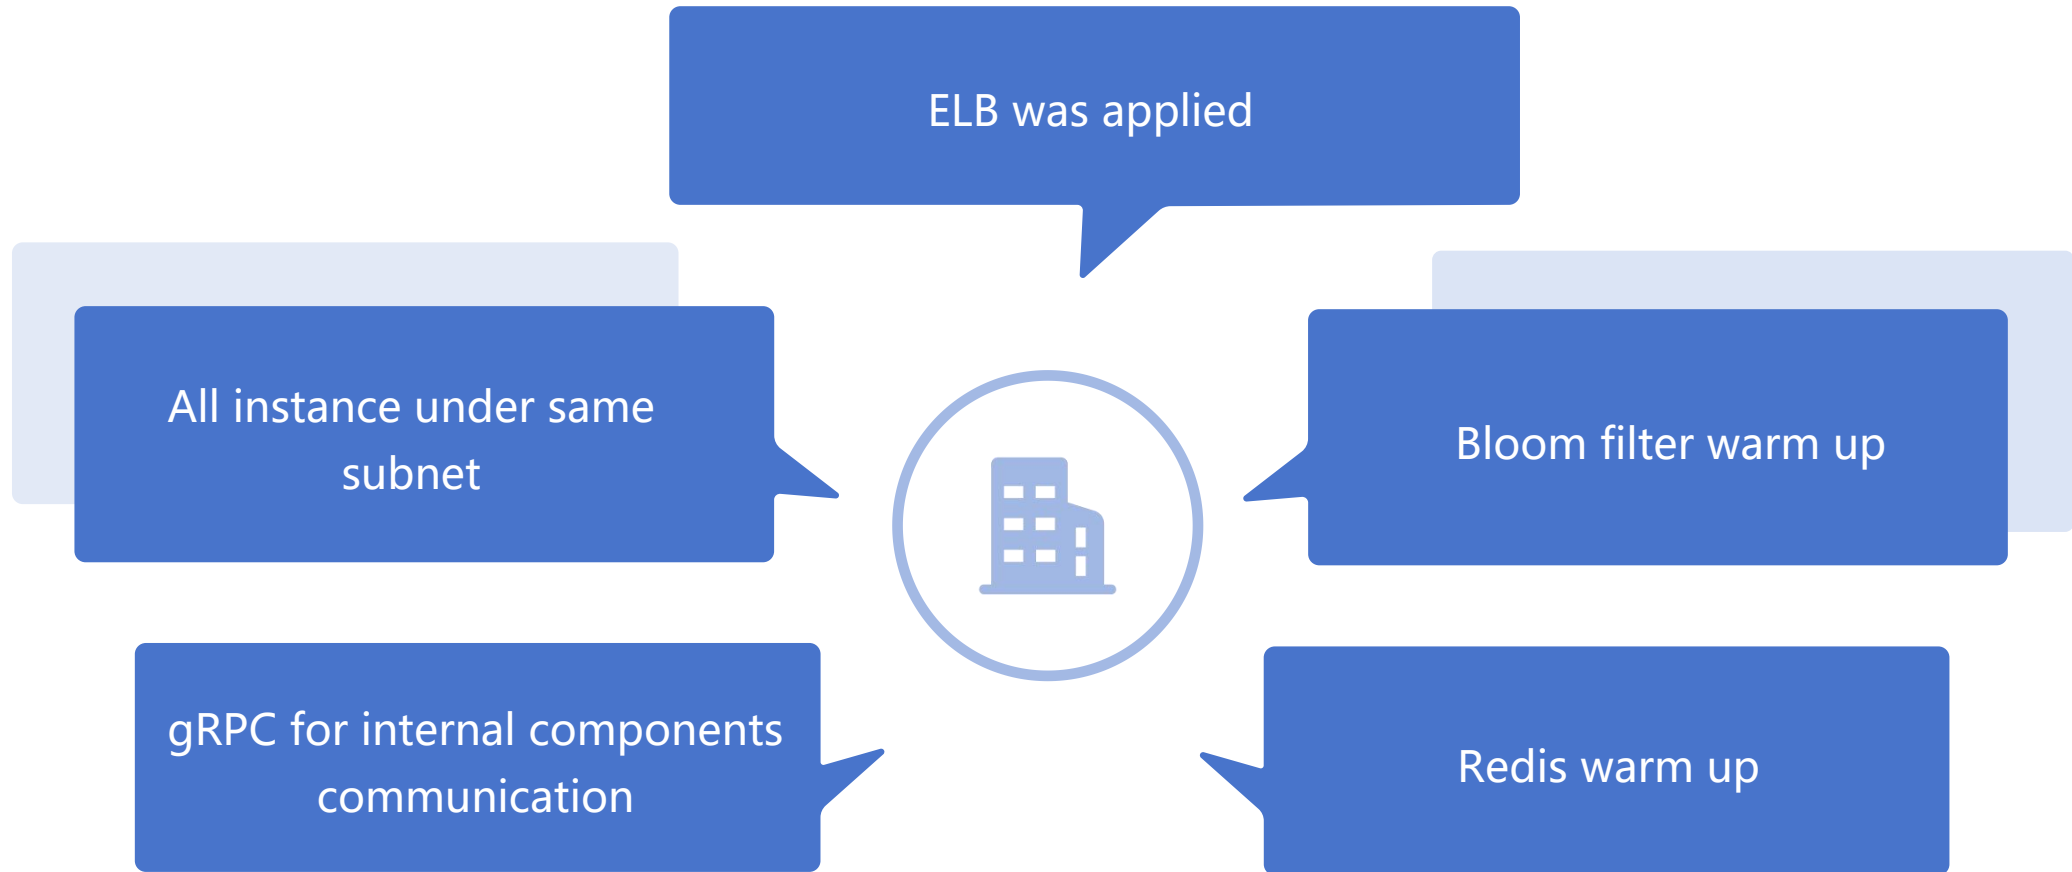
After:

- Faster read response times
- Higher throughput under concurrency
- Do not need to wait for cache update!

Why:



Some Tricks



Tune The System

01

Balance between availability and performance

02

We put all tunable parameters in config.properties file, easy to tune

03

Some major tunable parameters for writing

- Number of queues & Channels
- Message fetch amount & write batch size & write flush interval
- Write Service thread number

04

Some major tunable parameters for reading

- DB max pool size
- Read Service thread number
- Aggregation cache interval

Key Metrics Monitored

01 Write throughput

02 Read throughput

03 P99 latency

04 LRU hit amount

05 Cache hit amount

06 DB hit amount

07 Bloom filter negative amount

Result



Write

- Throughput: 10585.371017254156
- Mean response time: 92.447995 ms
- Min response time: 39.0 ms
- Max response time: 837.0 ms
- P99 response time: 321.0 ms



Read

- getAllVertical Throughput: 3795.52
- getSingleVertical Throughput: 4716.63
- getUniqueSkier Throughput: 4607.96

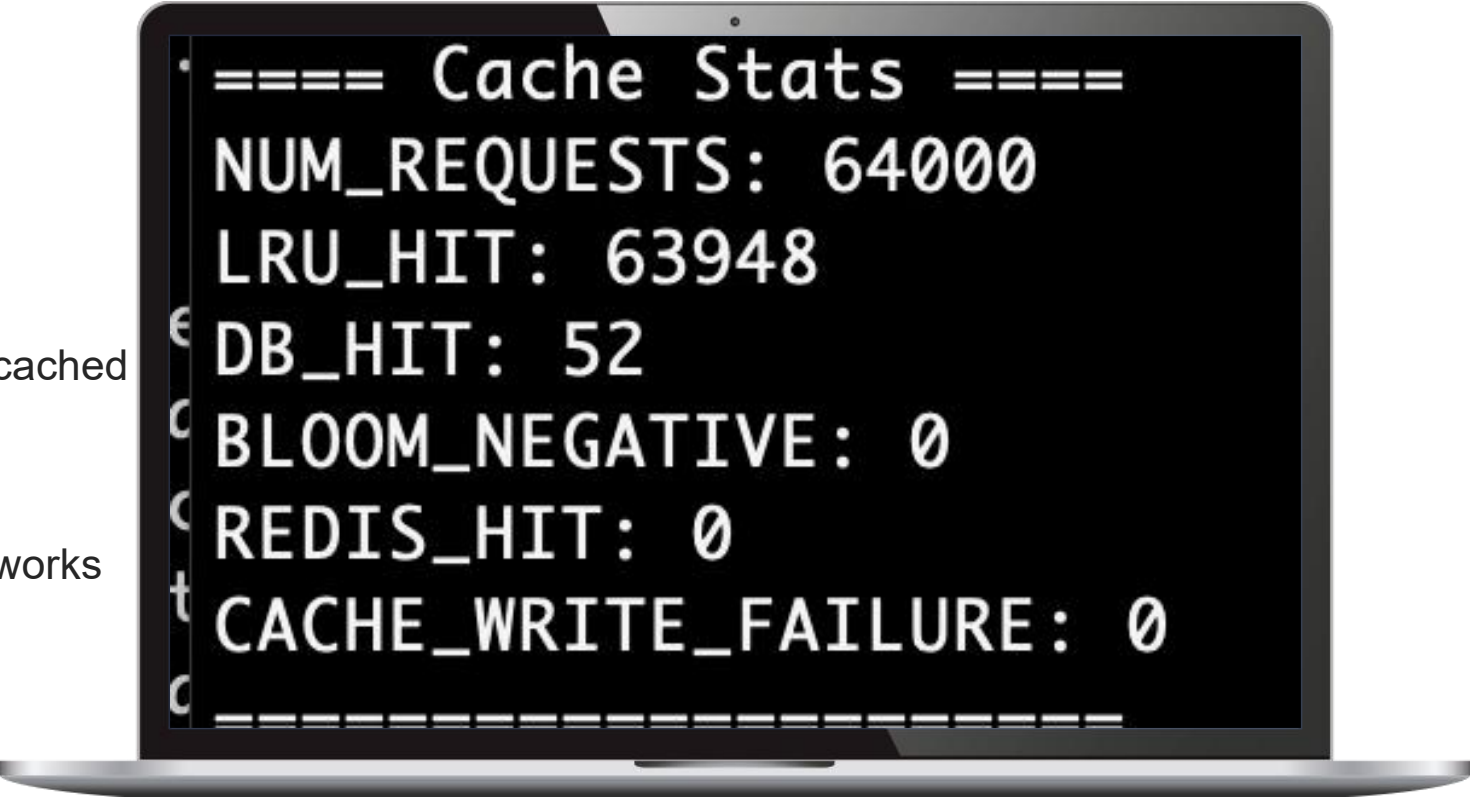
`GET/resorts/{resortID}/seasons/{seasonID}/day/{dayID}/skiers`

for getUniqueSkier api, most of read request were cached by LRU cache

10 resorts * 1 season * 3 days, total 30 keys, LRU works as expected!

Why still 52 DB hit?!

This is because under concurrent condition, there is a delay in writing to Redis and LRU cache between different thread.

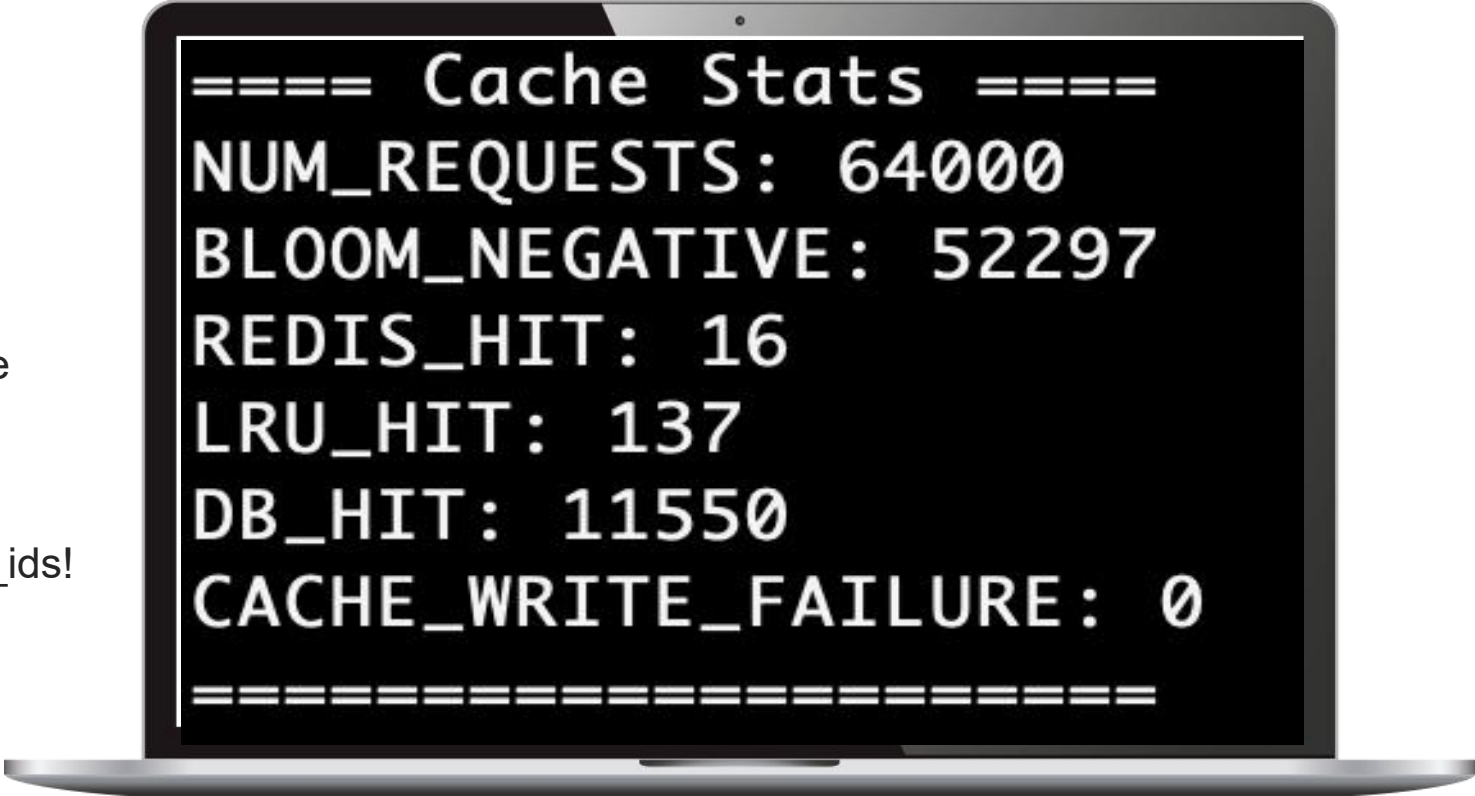
A laptop screen displaying cache statistics. The text is white on a black background, mimicking a terminal or command prompt. The statistics show a high number of requests (64000) with a very high hit rate (63948 LRU hits). There are 52 database hits, which is the focus of the text on the left. Bloom filter negatives, Redis hits, and cache write failures are all zero.

```
==== Cache Stats ====  
NUM_REQUESTS: 64000  
LRU_HIT: 63948  
DB_HIT: 52  
BLOOM_NEGATIVE: 0  
REDIS_HIT: 0  
CACHE_WRITE_FAILURE: 0  
=====
```

GET/skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID}

for getSingleVertical api, most of read request were filtered by bloom filters

This make sense because too much random skier_ids!

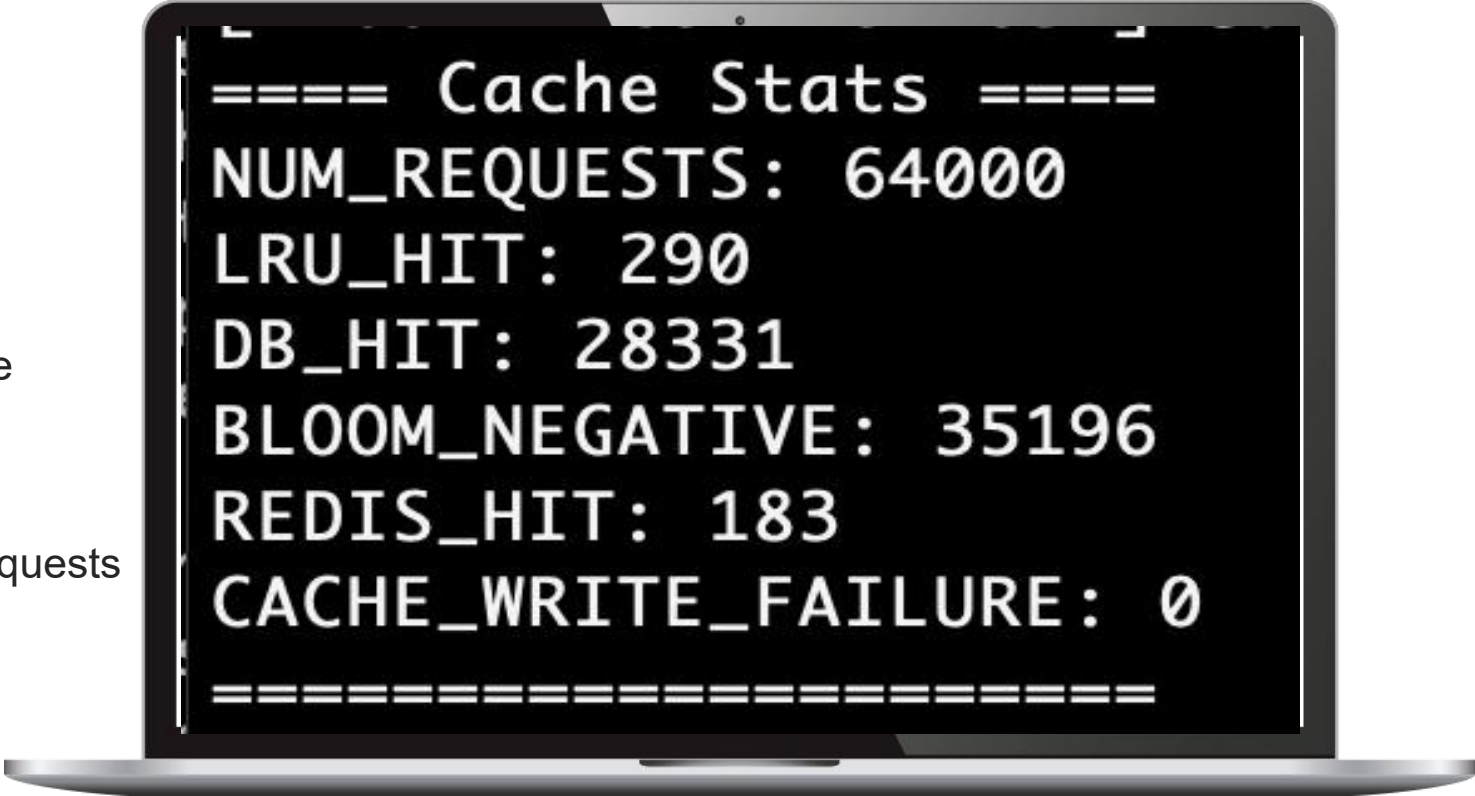


```
==== Cache Stats ====  
NUM_REQUESTS: 64000  
BLOOM_NEGATIVE: 52297  
REDIS_HIT: 16  
LRU_HIT: 137  
DB_HIT: 11550  
CACHE_WRITE_FAILURE: 0  
=====
```

GET/skiers/{skierID}/vertical

for getTotalVertical api, most of read request were filtered by bloom filters

Less random skier_ids than last api, but most of requests still filtered by bloom filters, as expected

A laptop screen with a black background and white text, displaying cache statistics. The text is centered and uses a monospaced font. The statistics are as follows:

```
==== Cache Stats ====
NUM_REQUESTS: 64000
LRU_HIT: 290
DB_HIT: 28331
BLOOM_NEGATIVE: 35196
REDIS_HIT: 183
CACHE_WRITE_FAILURE: 0
=====
```

Overall

Design Trade-off:
Availability & Performance > Consistency


The system is designed to be eventually
consistent

Stale data may exist during:

1. LRU cache refresh cycles
2. Redis cache write-back delays
3. Bloom filter update intervals

However, This trade-off gave us:

1. Faster read responses
2. Higher throughput under concurrent load
3. Better fault tolerance



Future Improvement

Distributed cache service (build a cluster)

Hybrid solution of multiple DB


Enable automatic recovery and back up

Maybe tried more powerful machines

Write-Behind Strategy for Cache Updates, write goes to DB first, then asyn flush to Cache

Infra buidlings: observability, data warehouse & data lake, ETL pipeline etc.





Thanks for watching Q&A